

Part 1: Classical Classification Algorithms

K-Nearest Neighbors (KNN)

- **Type:** Non-parametric, Instance-based (Lazy Learning).
- **Prediction Rule:** Assign the majority class among the k nearest neighbors (using a distance metric, typically Euclidean).
- **Pros**
 - Simple and easy to implement
 - No training phase (lazy learning)
 - Naturally handles non-linear decision boundaries
 - Effective for multi-class problems
- **Cons**
 - Slow inference: Computationally expensive at prediction time (calculates distance to all training points)
 - Memory intensive: Must store the entire training dataset
 - Sensitive to irrelevant or redundant features
 - Curse of dimensionality: Distance metrics lose meaning in very high dimensions; performance degrades with high-dimensional data
 - Requires careful choice of k, distance metric, and feature scaling

Bayes Optimal Classifier

- **Decision Rule:** $\hat{y} = \arg \max_y P(y|x)$
- **Pros**
 - Theoretically optimal (lowest possible error rate)
 - Provides true posterior probabilities
 - Serves as benchmark for all classifiers
- **Cons**
 - Intractable in practice (requires full knowledge of distributions)
 - Cannot be computed directly for most real problems

Naive Bayes (NB)

- **Type:** Generative Model ($P(x|y)$).
- **Decision Rule:** $\hat{y} = \arg \max_y P(y) \prod_i P(x_i|y)$ (using independence assumption)
- **Pros**
 - Extremely fast training and prediction
 - Performs surprisingly well even when independence assumption is violated
 - Works well with high-dimensional data (e.g., text/spam classification)
 - Robust to irrelevant features
 - Handles missing values naturally
 - Good for text classification (e.g., spam detection)
- **Cons**
 - Strong independence assumption: Assumes features are independent given the class (often violated)
 - Zero frequency problem: Requires smoothing (e.g., Laplace) for unseen features
 - Poor probability estimates (though good for classification; outputs often poorly calibrated and too extreme)

Linear Discriminant Analysis (LDA)

- **Type:** Generative Model.
- **Assumption:** Classes are Gaussian with shared covariance Σ .
- **Discriminant Function:** $\delta_k(x) = x^T \Sigma^{-1} \mu_k - \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log \pi_k$
- **Decision Boundary:** Linear (due to shared covariance).
- **Pros**
 - Supervised dimensionality reduction: Projects data to maximize class separability (max dimensions = $C - 1$); uses class labels
 - Optimal classifier if data is Gaussian with equal covariance matrices
 - Provides good class separation
 - Computationally efficient (closed-form solution)
- **Cons**
 - Strict assumptions: Requires Gaussian distribution and shared covariance matrix across classes
 - Sensitive to outliers
 - Limited to linear decision boundaries
 - Not suitable for non-Gaussian data

Naive Bayes vs. LDA

Feature	Naive Bayes	LDA
Assumption	Features are independent (Diagonal Covariance)	Features share correlations (Full Shared Covariance)
Similarities	Both assume Gaussian feature distributions (Gaussian NB); Both generative models; Both derive from Bayes theorem	
Differences	Faster and works better with limited data; Can have non-linear (quadratic) boundaries	Generally more accurate when correlations exist; Needs more data to estimate covariance matrix; Strictly linear boundaries
Data Efficiency	Better with small data	Needs more data to estimate covariance matrix

Logistic Regression (LR)

- **Type:** Discriminative Model ($P(y|x)$).
- **Model:** $P(y=1|x) = \sigma(w^T x + b) = \frac{1}{1+e^{-(w^T x + b)}}$ (Sigmoid function)
- **Pros**
 - Outputs well-calibrated probabilities (not just classifications)
 - Highly interpretable coefficients (feature importance)
 - Efficient training; easy to update with new data (using SGD)
 - Robust to noise
 - Works well when classes are linearly separable
- **Cons**
 - Inherently linear decision boundary (unless feature engineering)
 - Sensitive to outliers
 - Can overfit on high-dimensional data without regularization
 - Requires careful feature scaling

Logistic Regression vs. Generative Models (NB/LDA)

Feature	Logistic Regression	NB / LDA
Type Assumptions	Discriminative (models $P(y x)$ directly) Few assumptions about data distribution	Generative (models $P(x y)$ and $P(y)$) Strong assumptions (Gaussian, Independence)

Feature	Logistic Regression	NB / LDA
Similarities	All provide probabilistic outputs; All can handle multi-class problems	
Performance	Generally higher accuracy with sufficient data; Often performs better in practice; Doesn't assume feature distribution	Better with missing data or very small datasets

Support Vector Machines (SVM)

- **Type:** Discriminative (Max Margin).
- **Hard Margin Objective:** Maximize margin $\frac{2}{\|w\|}$ subject to $y_i(w^T x_i + b) \geq 1$
- **Pros**
 - Max margin: Theoretically robust to overfitting
 - Effective in high-dimensional spaces (even when $d > n$)
 - Memory efficient (defined only by support vectors)
 - Works well with clear margin separation
- **Cons**
 - No direct probability estimates (requires Platt scaling)
 - Sensitive to noise, parameter tuning (C), and choice
 - Poor performance with overlapping classes
 - Training computationally expensive on large datasets

Kernel SVM

- **Kernel Trick:** Maps data to higher dimension for non-linear boundaries.
- **Pros**
 - Handles non-linear data
 - Flexible with different kernels
 - Powerful for complex boundaries
- **Cons**
 - Computationally expensive ($O(n^2)$ to $O(n^3)$; kernel matrix)
 - Harder to interpret
 - Risk of overfitting with wrong kernel

Kernel SVM vs. Linear SVM

Feature	Linear SVM	Kernel SVM
Similarities	Both maximize margin; Both use support vectors	
Differences	Faster, simpler	Handles non-linear data but slower

Part 2: Regression & Optimization

Linear Regression Variants

1. **Standard Linear Regression**
 - Objective: Minimize $\sum(y_i - \hat{y}_i)^2$ (Sum of Squared Errors, SSE).
 - **Pros:** Simple and interpretable; Fast training/prediction; Closed-form solution; Good baseline
 - **Cons:** Assumes linear relationship; Sensitive to outliers; Unstable with correlated features/multicollinearity; Poor with non-linear data
2. **Ridge Regression (L2)**
 - Objective: Minimize $\sum(y_i - \hat{y}_i)^2 + \lambda \|w\|_2^2$
 - **Pros:** Handles multicollinearity; Reduces overfitting; Stable with noisy data
 - **Cons:** Doesn't perform feature selection; Bias introduced; Requires tuning lambda; Shrinks coefficients but keeps all

3. LASSO Regression (L1)

- Objective: Minimize $\sum(y_i - \hat{y}_i)^2 + \lambda\|w\|_1$
- **Pros:** Induces sparsity (sets some weights to zero); Acts as feature selection; Good for high-dimensional data; Handles sparse solutions
- **Cons:** Unstable with correlated features (can select only one from group)

Ridge vs. LASSO Regression

Feature	Ridge	LASSO
Similarities	Both regularized linear regression; Both shrink coefficients	
Differences	L2 penalty (shrinks but keeps all)	L1 penalty (can zero out coefficients)

Non-linear Regression

- **Pros**
 - Can model complex relationships
 - Flexible function forms
- **Cons**
 - Risk of overfitting
 - Harder to interpret
 - Computationally expensive

Kernel Ridge Regression

- **Pros**
 - Combines kernel trick with ridge
 - Handles non-linear data
 - Closed-form solution
- **Cons**
 - $O(n^3)$ complexity
 - No sparsity

Kernel Ridge vs. Ridge Regression

Feature	Ridge	Kernel Ridge
Similarities	Both use L2 regularization	
Differences	Linear only	Non-linear in input space

Support Vector Regression

- **Pros**
 - Robust to outliers (epsilon tube)
 - Handles non-linear via kernels
 - Good generalization
- **Cons**
 - Sensitive to parameters
 - Slow training on large data

Gradient Descent (GD) Variants

Variant	Batch Size	Pros	Cons
Batch GD	All Data	Stable convergence; exact gradient; Guaranteed convergence to minimum for convex functions	Slow on large datasets; High memory usage; Can get stuck in saddle/flat regions
Stochastic GD	Sample	Fast updates; Escapes local minima; Low memory; Good for online learning	Noisy convergence; May not converge exactly; Requires learning rate decay/tuning
Mini-Batch GD	N Samples	Best of both worlds: Balances speed and stability; Utilizes GPU vectorization; Lower variance than SGD; Good convergence properties	Requires tuning batch size; Still some noise in updates

Batch GD vs. Mini-batch GD vs. Stochastic GD

- Similarities:** All minimize the same objective function; All use gradient information
- Differences:** Batch: stable but slow; Mini-batch: compromise (most common in practice); Stochastic: fast but noisy

Coordinate Descent

- Pros**
 - Simple implementation
 - Fast for sparse problems
 - Effective for L1 regularization (Lasso)
- Cons**
 - May converge slowly
 - Order of coordinates matters
 - Not parallelizable easily

RANSAC (Random Sample Consensus)

- Purpose:** Robust fitting of models (e.g., lines, homographies) in the presence of many outliers.
- Pros:** Extremely robust to outliers/high outlier ratios; Simple concept
- Cons:** Non-deterministic; Requires threshold parameter and many iterations; No guarantee of optimal solution

Part 3: Unsupervised Learning & Clustering

Clustering Algorithms

- Exhaustive Clustering**
 - Pros:** Finds optimal clustering
 - Cons:** Computationally infeasible
- K-Means**
 - Objective: Minimize $\sum_{k=1}^K \sum_{i \in C_k} \|x_i - \mu_k\|^2$ (Within-cluster sum of squares)
 - Hard assignment; Assumes spherical clusters of similar size.
 - Pros:** Simple and fast; Scales well; Easy to implement
 - Cons:** Sensitive to initialization and outliers; Requires choosing k
- Gaussian Mixture Models (GMM)**
 - Probability: $p(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x|\mu_k, \Sigma_k)$
 - Soft assignment (probabilistic); Assumes elliptical clusters; Uses Expectation-Maximization (EM).
 - Pros:** Handles elliptical clusters; Can model complex distributions; Provides probabilities
 - Cons:** Sensitive to initialization; Can overfit; Computationally heavier; Slower

K-Means vs. Gaussian Mixture Model

Feature	K-Means	GMM
Similarities	Both centroid-based; Use EM-like optimization	
Differences	Hard assignment, spherical clusters	Soft assignment, elliptical clusters

Expectation-Maximization (EM) Algorithm

- Pros
 - Handles latent variables
 - Monotonic improvement
 - Flexible for many models
- Cons
 - Local optima
 - Slow convergence
 - Sensitive to initialization

Part 4: Dimensionality Reduction & Feature Processing

Dimensionality Reduction

- Pros
 - Reduces computation
 - Removes noise/redundancy
 - Helps visualization
 - Mitigates curse of dimensionality
- Cons
 - Loss of information
 - Harder interpretation

Feature Selection

- Pros
 - Maintains interpretability
 - Reduces overfitting
 - Faster training
- Cons
 - May miss combined effects
 - Computationally expensive (some methods)

Dimensionality Reduction vs. Feature Selection

Feature	Feature Selection	Dimensionality Reduction
Similarities	Both reduce features	
Differences	Subset of original features	New transformed features

Linear Dimensionality Reduction

- Pros
 - Fast and simple
 - Interpretable
- Cons
 - Assumes linear relationships
 - May miss complex structures

Singular Value Decomposition (SVD)

- **Pros**
 - Numerically stable
 - Basis for many methods
 - Handles any matrix
- **Cons**
 - Linear only
 - Computationally intensive for large matrices
- The mathematical matrix factorization technique that underpins PCA; More numerically stable than eigendecomposition.

Principal Component Analysis (PCA)

- **Unsupervised.**
- Finds orthogonal directions of maximum variance (eigenvectors of covariance matrix, ordered by eigenvalues).
- **Pros**
 - Maximizes variance
 - Orthogonal components
 - Reduces dimensions effectively
- **Cons**
 - Sensitive to scaling
 - Assumes linear correlations
 - Interprets variance as importance (“Variance” does not always equal “Information”)

PCA vs. SVD vs. Dimensionality Reduction

- **Similarities:** PCA uses SVD; All linear techniques
- **Differences:** SVD: general matrix decomposition; PCA: SVD on centered data for variance maximization; Dimensionality reduction: broader category

Feature Normalization / Data Whitening

- **Pros**
 - Improves convergence speed/optimization
 - Prevents feature dominance
 - Essential for distance-based algorithms
 - Decorrelates features; Equalizes variances
- **Cons**
 - Can leak test information if not careful
 - May distort data distribution/amplify noise
 - Assumes linear correlations
 - Computationally expensive (whitening)

Output Transformation

- **Pros**
 - Can stabilize variance
 - Make data more normal-like
 - Help meet model assumptions
- **Cons**
 - Harder interpretation
 - Need inverse transform for predictions

Part 5: Deep Learning

Architectures

- **Multi-Layer Perceptron (MLP)**

- Dense connections; Good for tabular data.

- **Pros**

- * Universal approximator
 - * Handles non-linear data
 - * Flexible architecture

- **Cons**

- * Black box
 - * Prone to overfitting on images
 - * Requires much data

- **Convolutional Neural Network (CNN)**

- **Inductive Bias:** Translation invariance and locality; Parameter sharing.

- **Pros**

- * Excellent for images/grid data
 - * Learns hierarchical features
 - * Efficient due to parameter sharing

- **Cons**

- * Requires large data
 - * Computationally expensive
 - * Less effective on non-spatial data

MLP vs. CNN

Feature	MLP	CNN
Similarities	Both deep neural networks; Both use backpropagation	
Differences	Fully connected	Local connectivity, weight sharing

Regularization & Training Techniques

- **Dropout**

- Randomly deactivates neurons during training; Prevents co-adaptation (ensemble effect).

- **Pros**

- * Effective regularization
 - * Reduces co-adaptation

- **Cons**

- * Increases training time
 - * Requires larger networks
 - * Inference needs scaling

- **Early Stopping**

- Stops training when validation error rises.

- **Pros**

- * Prevents overfitting
 - * Saves training time
 - * Simple to implement

- **Cons**

- * Requires validation set

- * May stop too early
- * Sensitive to patience parameter

- **Data Augmentation**

- Artificially expands dataset (flip, rotate, crop).
- **Pros**

- * Increases effective dataset size
- * Improves generalization
- * Reduces overfitting
- * Cost-effective; Crucial for computer vision

- **Cons**

- * Can introduce unrealistic samples
- * Increases training time
- * Domain-specific design needed

- **Batch Normalization / Whitening**

- Normalizes layer inputs (zero mean, unit variance).
- **Pros**
- * Stabilizes learning
- * Allows higher learning rates
- **Cons** (see Data Whitening above)

Early Stopping vs. Dropout

Feature	Early Stopping	Dropout
Similarities	Both prevent overfitting; Both regularization techniques	
Differences	Stops training early	Random neuron drop during training

Weight Initialization

- **Methods**

- Zero (bad)
- Random normal/uniform
- Xavier/Glorot (for sigmoid/tanh)
- He (for ReLU)
- Orthogonal

Hyperparameter Tuning

- **Procedures**

- Grid search
- Random search
- Bayesian optimization
- Manual tuning
- Cross-validation

Learning Curve Interpretation

- **Accuracy still going up:** Model has capacity to learn more → Needs more data or longer training
- **Huge gap between train and validation accuracy:** High variance/overfitting → Needs regularization or more data
- **Little gap between train and validation:** Low variance → If both low: high bias/underfitting

Part 6: Computer Vision & Image Processing

Tasks

- **Object Detection:** Locates and classifies objects with bounding boxes (e.g., YOLO, Faster R-CNN)
- **Semantic Segmentation:** Pixel-level classification (all instances of same class share label)
- **Instance Segmentation:** Pixel-level classification distinguishing individual objects (different instances of same class)

Image Restoration & Processing

- **Denoising:** Removing noise (e.g., Filtering, wavelet, deep learning like DnCNN)
 - **Pros:** Improves image quality; Enhances downstream tasks
 - **Cons:** May blur details; Computationally intensive (deep methods)
- **Deblurring:** Removes motion/camera shake blur (harder than denoising)
- **Super-Resolution:** Upscaling images; GANs often used to hallucinate details
 - **Pros:** Enhances details
 - **Cons:** Can introduce artifacts
- **Colorization:** Adds color to grayscale
 - **Pros:** Enhances old photos
 - **Cons:** Subjective; Can be inaccurate
- **Compression:** Reduces file size
 - **Pros:** Storage/transmission efficiency
 - **Cons:** Loss of quality (lossy)
- **Fusion:** Combines multiple images (multi-modal)
 - **Pros:** Richer information
 - **Cons:** Alignment issues

Quality Metrics

- **MSE/PSNR**
 - **Pros:** Simple, differentiable; Good for optimization
 - **Cons:** Poor correlation with human perception (penalizes slight shifts heavily); Sensitive to outliers
- **SSIM (Structural Similarity)**
 - Captures luminance, contrast, and structure.
 - **Pros:** Closer to human vision; Better perceptual quality
 - **Cons:** More complex; Sensitive to scaling/alignment
- **LPIPS (Learned Perceptual Image Patch Similarity)**
 - Uses deep network features.
 - **Pros:** Best match for human perception
 - **Cons:** Computationally heavy; Requires pretrained network

MSE vs. SSIM

Feature	MSE	SSIM
Similarities	Both full-reference metrics	
Differences	Pixel-wise error	Structural/perceptual similarity

General Notes

- **Model Comparison:** Use multiple metrics (accuracy, speed, interpretability) and validation techniques